

特 別 研 究 報 告

題 目

This Is
My Bachelor Thesis

指 導 教 員

Professor Jane Doe
Associate Professor Joe Smith

報 告 者

John Doe

平成 28 年 2 月 吉日

大阪大学基礎工学部システム科学科
知能システム学コース

Abstract

This is abstract.

Contents

Abstract	i
1 Introduction	1
1.1 First of All	1
1.2 Next of First of All	1
1.2.1 First Subsection	1
2 2nd Chapter	2
2.1 Markov Decision Processes and Reinforcement Learning	2
2.2 Stochastic Discrete Event Systems	4
2.3 Linear Temporal Logic and Automata	5
3 3rd Chapter	8
3.1 Reinforcement-Learning-Based Synthesis of Control Policy	8
3.2 Example	10
4 4th Chapter	14
5 Conclusions	15
A Proofs	16
Acknowledgment	19
References	20

Chapter 1

Introduction

Test of Theorem.

1.1 First of All

Theorem 1.1 (First Theorem) This is Theorem.

Lemma 1.1 (First Lemma) This is Lemma.

Proposition 1.1 (First Proposition) This is Proposition.

Corollary 1.1 (First Corollary) This is Corollary.

Proof This is true, obviously.

□

Definition 1.1 (First Definition) We call this “That”.

Example 1.1 (First Example) This is an example.

1.2 Next of First of All

1.2.1 First Subsection

First Subsubsection

Test of Equation.

$$a = b + c \tag{1.1}$$

Table 1.1 Test of Table Caption

Fig. 1.1 Test of Figure Caption

Chapter 2

2nd Chapter

2.1 Markov Decision Processes and Reinforcement Learning

Definition 2.1 (Labeled Markov Decision Process) A (labeled) Markov decision process (MDP) is a tuple $M = (S, A, \mathcal{A}, P, s_{init}, AP, L)$, where S is a finite set of states, A is a finite set of actions, $\mathcal{A} : S \rightarrow 2^A$ is a mapping that maps each state to the set of possible actions at the state, $P : S \times S \times A \rightarrow [0, 1]$ is a transition probability such that $\sum_{s' \in S} P(s'|s, a) = 1$ for any state $s \in S$ and any action $a \in \mathcal{A}(s)$, $s_{init} \in S$ is the initial state, AP is a finite set of atomic propositions, and $L : S \times A \times S \rightarrow 2^{AP}$ is a labeling function that assigns a set of atomic propositions to each transition $(s, a, s') \in S \times A \times S$.

In the MDP M , an infinite path starting from a state $s_0 \in S$ is defined as a sequence $\rho = s_0 a_0 s_1 \dots \in S(AS)^\omega$ such that $P(s_{i+1}|s_i, a_i) > 0$ for any $i \in \mathbb{N}_0$, where \mathbb{N}_0 is the set of natural numbers including zero. A finite path is a finite sequence in $S(AS)^*$. In addition, we sometimes represent ρ as ρ_{init} to emphasize that ρ starts from $s_0 = s_{init}$. For a path $\rho = s_0 a_0 s_1 \dots$, we define the corresponding labeled path $L(\rho) = L(s_0, a_0, s_1) L(s_1, a_1, s_2) \dots \in (2^{AP})^\omega$. $InfPath^M$ (resp., $FinPath^M$) is defined as the set of infinite (resp., finite) paths starting from $s_0 = s_{init}$ in the MDP M . For each finite path ρ , $last(\rho)$ denotes its last state.

Definition 2.2 (Policy) A policy on an MDP M is defined as a mapping $\pi : FinPath^M \times \mathcal{A}(last(\rho)) \rightarrow [0, 1]$. A policy π is a *positional* policy if for any $\rho \in FinPath^M$ and any $a \in \mathcal{A}(last(\rho))$, it holds that $\pi(\rho, a) = \pi(last(\rho), a)$ and there exists $a' \in \mathcal{A}(last(\rho))$ such that

$$\pi(\rho, a) = \begin{cases} 1 & \text{if } a = a', \\ 0 & \text{otherwise.} \end{cases}$$

Let $InfPath_\pi^M$ (resp., $FinPath_\pi^M$) be the set of infinite (resp., finite) paths starting

from $s_0 = s_{init}$ in the MDP M under a policy π . The behavior of an MDP M under a policy π is defined on a probability space $(InfPath_\pi^M, \mathcal{F}_{InfPath_\pi^M}, Pr_\pi^M)$.

Definition 2.3 (Markov chain) A Markov chain induced by an MDP M with a positional policy π is a tuple $MC_\pi = (S_\pi, P_\pi, s_0, AP, L)$, where $S_\pi = S$, $P_\pi(s'|s) = P(s'|s, a)$ for $s, s' \in S$ and $a \in \mathcal{A}(s)$ such that $\pi(s, a) = 1$. The state set S_π of MC_π can be represented as a disjoint union of a set of transient states T_π and closed irreducible sets of recurrent states R_π^j with $j \in \{1, \dots, h\}$, as $S_\pi = T_\pi \sqcup R_\pi^1 \sqcup \dots \sqcup R_\pi^h$ [18]. In the following, we say a “recurrent class” instead of a “closed irreducible set of recurrent states” for simplicity.

In an MDP M , we define a reward function $R : S \times A \times S \rightarrow \mathbb{R}$, where \mathbb{R} is the set of real numbers. The function denotes the immediate scalar bounded reward received after the agent performs an action a at a state s and reaches a next state s' as a result.

Definition 2.4 (Expected discounted reward for MDPs) For a policy π on an MDP M , any state $s \in S$, and a reward function R , we define the expected discounted reward as

$$V^\pi(s) = \mathbb{E}^\pi \left[\sum_{n=0}^{\infty} \gamma^n R(S_n, A_n, S_{n+1}) | S_0 = s \right],$$

where \mathbb{E}^π denotes the expected value given that the agent follows the policy π from the state s and $\gamma \in [0, 1)$ is a discount factor. Intuitively, the magnitude of the discount factor γ determines how much we consider rewards received in the future. The function $V^\pi(s)$ is often referred to as a state-value function under the policy π . For any state-action pair $(s, a) \in S \times A$, we define an action-value function $Q^\pi(s, a)$ under the policy π as follows.

$$Q^\pi(s, a) = \mathbb{E}^\pi \left[\sum_{n=0}^{\infty} \gamma^n R(S_n, A_n, S_{n+1}) | S_0 = s, A_0 = a \right].$$

We have the following recursively equation for the state-value function and the action-value function.

$$\begin{aligned} V^\pi(s) &= \mathbb{E}^\pi \left[\sum_{n=0}^{\infty} \gamma^n R(S_n, A_n, S_{n+1}) | S_0 = s \right] \\ &= \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in S} P(s'|s, a) \mathbb{E}^\pi \left[\sum_{n=0}^{\infty} \gamma^n R(S_n, A_n, S_{n+1}) | S_0 = s, A_0 = a, S_1 = s' \right] \\ &= \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in S} P(s'|s, a) \{ R(s, a, s') + \gamma \mathbb{E}^\pi \left[\sum_{n=0}^{\infty} \gamma^n R(S_n, A_n, S_{n+1}) | S_1 = s' \right] \} \end{aligned}$$

$$= \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma V^\pi(s')\},$$

by the definition of the action-value function, it holds That

$$\begin{aligned} Q^\pi(s, a) &= \max_{a \in \mathcal{A}(s)} V^\pi(s) \\ &= \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma V^\pi(s')\} \\ &= \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma \sum_{a' \in \mathcal{A}(s')} \pi(s', a') Q^\pi(s', a')\}. \end{aligned}$$

The above equations are called the *Bellman equation*.

Definition 2.5 (Optimal policy) For any state s in S , a policy π^* is optimal if

$$\pi^* \in \arg \max_{\pi \in \Pi^{pos}} V^\pi(s),$$

where Π^{pos} is the set of positional policies over the state set S .

We have the following *Bellman optimality functions* by the optimal policy.

$$\begin{aligned} V^{\pi^*}(s) &= \max_{\pi \in \Pi^{pos}} V^\pi(s) \\ &= \max_{\pi \in \Pi^{pos}} \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma V^\pi(s')\} \\ &= \max_{a \in \mathcal{A}(s)} \left[\sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma V^{\pi^*}(s')\} \right], \end{aligned}$$

$$\begin{aligned} Q^{\pi^*}(s, a) &= \max_{\pi \in \Pi^{pos}} Q^\pi(s, a) \\ &= \max_{\pi \in \Pi^{pos}} \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma \sum_{a' \in \mathcal{A}(s')} \pi(s', a') Q^\pi(s', a')\} \\ &= \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma \max_{\pi \in \Pi^{pos}} \sum_{a' \in \mathcal{A}(s')} \pi(s', a') Q^\pi(s', a')\} \\ &= \sum_{s' \in S} P(s'|s, a) \{R(s, a, s') + \gamma \max_{a' \in \mathcal{A}(s')} Q^{\pi^*}(s', a')\}. \end{aligned}$$

2.2 Stochastic Discrete Event Systems

Definition 2.6 (Stochastic discrete event system) We represent a stochastic discrete event system (DES) as a labeled Markov decision process (MDP). A DES is a tuple

$D = (S, E, \mathcal{E}, P_T, P_E, s_{init}, AP, L)$, where S is a finite set of states; E is a finite set of events; $\mathcal{E} : S \rightarrow E$ is a mapping that maps each state to the set of feasible events at the state; $P_T : S \times S \times E \rightarrow [0, 1]$ is a transition probability such that $\sum_{s' \in S} P_T(s'|s, e) = 1$ for any state $s \in S$ and any event $e \in \mathcal{E}(s)$ and $P_T(s'|s, e) = 0$ for any $e \notin \mathcal{E}(s)$; $P_E : E \times S \times 2^E \rightarrow [0, 1]$ is the probability that an event occurs under a subset $\pi \in \mathcal{E}(s)$ of events allowed to occur at the state $s \in S$ such that $\sum_{e \in \pi} P_E(e|s, \pi) = 1$ and we call the subset the control pattern; for any $(s', s, \pi) \in S \times S \times 2^E$, we define the probability $P : S \times S \times 2^E \rightarrow [0, 1]$ such that $P(s'|s, \pi) = \sum_{e \in \pi} P_E(e|s, \pi)P_T(s'|s, e)$ and $\sum_{s' \in S} P(s'|s, \pi) = 1$; $s_{init} \in S$ is the initial state; AP is a finite set of atomic propositions; and $L : S \times E \times S \rightarrow 2^{AP}$ is a labeling function that assigns a set of atomic propositions to each transition $(s, e, s') \in S \times E \times S$. We assume that E can be partitioned into the set of controllable events E_c and the set of uncontrollable events E_{uc} such that $E_c \cup E_{uc} = E$ and $E_c \cap E_{uc} = \emptyset$. Note that each event e occurs probabilistically depending on only the current state and the subset of feasible events at the state given by a controller.

In the DES D , an infinite path for the DES starting from a state $s_0 \in S$ is defined as a sequence $\rho^D = s_0 \pi_0 e_0 s_1 \dots \in S(2^E E S)^\omega$ such that $P_E(e_i|s_i, \pi_i) > 0$ and $P_T(s_{i+1}|s_i, e_i) > 0$ for any $i \in \mathbb{N}_0$. A finite path for the DES is a finite sequence in $S(2^E E S)^*$. In addition, we sometimes represent ρ^D as ρ_{init}^D to emphasize that ρ^D starts from $s_0 = s_{init}$. For a path $\rho^D = s_0 \pi_0 e_0 s_1 \dots$, we define the corresponding labeled path $L(\rho^D) = L(s_0, e_0, s_1)L(s_1, e_1, s_2) \dots \in (2^{AP})^\omega$. For simplicity, we often represent infinite (resp., finite) path for the DES as infinite (resp., finite) path and omit superscript D of the paths. $InfPath^D$ (resp., $FinPath^D$) is defined as the set of infinite (resp., finite) paths starting from $s_0 = s_{init}$ in the DES D . For each finite path ρ , $last(\rho)$ denotes its last state.

We define the supervisor as a controller for the DES that restricts the behaviors of the DES to satisfy a given specification.

Definition 2.7 For the DES M , a supervisor $SV : FinPath^M \rightarrow 2^E$ is defined as a mapping that maps each finite path to a set of allowed events at the finite path and we call the set the control pattern. In the following, the supervisor we consider is *state-based*, namely for any $\rho \in FinPath^M$, $SV(\rho) = SV(last(\rho))$. Note that the relation $E_{uc} \subset SV(\rho) \subset E$ holds for any $\rho \in FinPath^M$.

2.3 Linear Temporal Logic and Automata

In our proposed method, we use linear temporal logic (LTL) formulas to describe various constraints or properties and to systematically assign corresponding rewards. LTL formulas are constructed from a set of atomic propositions, Boolean operators, and temporal operators. We use the standard notations for the Boolean operators: \top (true), \neg

(negation), and \wedge (conjunction). LTL formulas over a set of atomic propositions AP are recursively defined as

$$\varphi ::= \top \mid \alpha \in AP \mid \varphi_1 \wedge \varphi_2 \mid \neg\varphi \mid \mathbf{X}\varphi \mid \varphi_1 \mathbf{U}\varphi_2,$$

where φ , φ_1 , and φ_2 are LTL formulas. Additional Boolean operators are defined as $\perp := \neg\top$, $\varphi_1 \vee \varphi_2 := \neg(\neg\varphi_1 \wedge \neg\varphi_2)$, and $\varphi_1 \Rightarrow \varphi_2 := \neg\varphi_1 \vee \varphi_2$. The operators \mathbf{X} and \mathbf{U} are called “next” and “until”, respectively. Using the operator \mathbf{U} , we define two temporal operators: (1) *eventually*, $\mathbf{F}\varphi := \top \mathbf{U}\varphi$ and (2) *always*, $\mathbf{G}\varphi := \neg\mathbf{F}\neg\varphi$.

Let M be an MDP. For an infinite path $\rho = s_0 a_0 s_1 \dots$ of M with $s_0 \in S$, let $\rho[i]$ be the i -th state of ρ i.e., $\rho[i] = s_i$ and let $\rho[i:]$ be the i -th suffix $\rho[i:] = s_i a_i s_{i+1} \dots$.

Definition 2.8 (LTL semantics) For an LTL formula φ , an MDP M , and an infinite path $\rho = s_0 a_0 s_1 \dots$ of M with $s_0 \in S$, the satisfaction relation $M, \rho \models \varphi$ is recursively defined as follows.

$$\begin{aligned} M, \rho &\models \top, \\ M, \rho &\models \alpha \in AP \Leftrightarrow \alpha \in L(s_0, a_0, s_1), \\ M, \rho &\models \varphi_1 \wedge \varphi_2 \Leftrightarrow M, \rho \models \varphi_1 \wedge M, \rho \models \varphi_2, \\ M, \rho &\models \neg\varphi \Leftrightarrow M, \rho \not\models \varphi, \\ M, \rho &\models \mathbf{X}\varphi \Leftrightarrow M, \rho[1:] \models \varphi, \\ M, \rho &\models \varphi_1 \mathbf{U}\varphi_2 \Leftrightarrow \exists j \geq 0, M, \rho[j:] \models \varphi_2 \wedge \forall i, 0 \leq i < j, M, \rho[i:] \models \varphi_1. \end{aligned}$$

The next operator \mathbf{X} requires that φ is satisfied by the next state suffix of ρ . The until operator \mathbf{U} requires that φ_1 holds true until φ_2 becomes true over the path ρ . In the following, we write $\rho \models \varphi$ for simplicity without referring to MDP M .

For any policy π , we denote the probability of all paths starting from s_{init} on the MDP M that satisfy an LTL formula φ under the policy π as

$$Pr_{\pi}^M(s_{init} \models \varphi) := Pr_{\pi}^M(\{\rho_{init} \in InfPath_{\pi}^M; \rho_{init} \models \varphi\}).$$

We say that an LTL formula φ is satisfied by a positional policy π if

$$Pr_{\pi}^M(s_{init} \models \varphi) > 0.$$

Any LTL formula φ can be converted into various automata, namely finite state machines that recognize all words satisfying φ . We define a generalized Büchi automaton at the beginning, and then introduce a limit-deterministic Büchi automaton.

Definition 2.9 (Transition-based generalized Büchi automaton) A transition-based generalized Büchi automaton (tGBA) is a tuple $B = (X, x_{init}, \Sigma, \delta, \mathcal{F})$, where X is a finite set of states, $x_{init} \in X$ is the initial state, Σ is an input alphabet, $\delta \subset X \times \Sigma \times X$

is a set of transitions, and $\mathcal{F} = \{F_1, \dots, F_n\}$ is an acceptance condition, where for each $j \in \{1, \dots, n\}$, $F_j \subset \delta$ is a set of accepting transitions and called an accepting set.

Let Σ^ω be the set of all infinite words over Σ and let an infinite run be an infinite sequence $r = x_0\sigma_0x_1 \dots \in X(\Sigma X)^\omega$ where $(x_i, \sigma_i, x_{i+1}) \in \delta$ for any $i \in \mathbb{N}_0$. An infinite word $w = \sigma_0\sigma_1 \dots \in \Sigma^\omega$ is accepted by B_φ if and only if there exists an infinite run $r = x_0\sigma_0x_1 \dots$ starting from $x_0 = x_{init}$ such that $inf(r) \cap F_j \neq \emptyset$ for each $F_j \in \mathcal{F}$, where $inf(r)$ is the set of transitions that occur infinitely often in the run r .

Definition 2.10 (Sink state) A sink state in state set X of an augmented tLDBA $\bar{B}_\varphi = (\bar{X}, \bar{x}_{init}, \bar{\Sigma}, \bar{\delta}, \bar{\mathcal{F}})$ is defined as a state such that there exist no accepting transition of \bar{B}_φ that is accessible from the state. We denote the set of sink states as *SinkSet*.

Definition 2.11 (Transition-based limit-deterministic generalized Büchi automaton) A tGBA $B = (X, x_{init}, \Sigma, \delta, \mathcal{F})$ is limit-deterministic (tLDBA) if the following conditions hold.

- $\exists X_{initial}, X_{final} \subset X$ s.t. $X = X_{initial} \cup X_{final} \wedge X_{initial} \cap X_{final} = \emptyset$,
- $F_j \subset X_{final} \times \Sigma \times X_{final}, \forall j \in \{1, \dots, n\}$,
- $|\{(x, \sigma, x') \in \delta; x' \in X_{initial}\}| \leq 1, \forall x \in X_{initial}, \forall \sigma \in \Sigma$,
- $|\{(x, \sigma, x') \in \delta; x' \in X_{final}\}| \leq 1, \forall x \in X_{final}, \forall \sigma \in \Sigma$,
- $|\{(x, \sigma, x') \in \delta; x' \in X_{initial}\}| = 0, \forall x \in X_{final}, \forall \sigma \in \Sigma$.

A tLDBA is a tGBA whose state set can be partitioned into the initial part $X_{initial}$ and the final part X_{final} , and they are connected by a single “guess”. The final part has all accepting sets. The transitions in each part are deterministic. It is known that, for any LTL formula φ , there exists a tLDBA that accepts all words satisfying φ [13]. In particular, we represent a tLGBA recognizing an LTL formula φ as B_φ , whose input alphabet is given by $\Sigma = 2^{AP}$.

Chapter 3

3rd Chapter

3.1 Reinforcement-Learning-Based Synthesis of Control Policy

We introduce an automaton augmented with binary vectors. The automaton can explicitly represent whether transitions in each accepting set occur at least once, and ensure transitions in each accepting set occur infinitely often.

Let $V = \{(v_1, \dots, v_n)^T ; v_i \in \{0, 1\}, i \in \{1, \dots, n\}\}$ be a set of binary-valued vectors, and let $\mathbf{1}$ and $\mathbf{0}$ be the n -dimensional vectors with all elements 1 and 0, respectively. In order to augment a tLDBA B_φ , we introduce three functions $visitf : \delta \rightarrow V$, $reset : V \rightarrow V$, and $Max : V \times V \rightarrow V$ as follows. For any $e \in \delta$, $visitf(e) = (v_1, \dots, v_n)^T$, where

$$v_i = \begin{cases} 1 & \text{if } e \in F_i, \\ 0 & \text{otherwise.} \end{cases}$$

For any $v \in V$,

$$reset(v) = \begin{cases} \mathbf{0} & \text{if } v = \mathbf{1}, \\ v & \text{otherwise.} \end{cases}$$

For any $v, u \in V$, $Max(v, u) = (l_1, \dots, l_n)^T$, where $l_i = \max\{v_i, u_i\}$ for any $i \in \{1, \dots, n\}$.

Intuitively, each vector v represents which accepting sets have been visited. The function $visitf$ returns a binary vector whose i -th element is 1 if and only if a transition in the accepting set F_i occurs. The function $reset$ returns the zero vector $\mathbf{0}$ if at least one transition in each accepting set has occurred after the latest reset. Otherwise, it returns the input vector without change.

Definition 3.1 For a tLDBA $B_\varphi = (X, x_{init}, \Sigma, \delta, \mathcal{F})$, its augmented automaton is a tLDBA $\bar{B}_\varphi = (\bar{X}, \bar{x}_{init}, \bar{\Sigma}, \bar{\delta}, \bar{\mathcal{F}})$, where $\bar{X} = X \times V$, $\bar{x}_{init} = (x_{init}, \mathbf{0})$, $\bar{\Sigma} = \Sigma$, $\bar{\delta}$ is defined as $\bar{\delta} = \{((x, v), \bar{\sigma}, (x', v')) \in \bar{X} \times \bar{\Sigma} \times \bar{X} ; (x, \bar{\sigma}, x') \in$

$\delta, v' = \text{reset}(\text{Max}(v, \text{visitf}((x, \bar{\sigma}, x'))))$, and $\bar{\mathcal{F}} = \{\bar{F}_1, \dots, \bar{F}_n\}$ is defined as $\bar{F}_i = \{((x, v), \bar{\sigma}, (x', v')) \in \bar{\delta} ; (x, \sigma, x') \in F_i, v_i = 0, \text{visitf}((x, \bar{\sigma}, x'))_i = 1\}$ for each $i \in \{1, \dots, n\}$, where $\text{visitf}((x, \bar{\sigma}, x'))_i$ is the i -th element of $\text{visitf}((x, \bar{\sigma}, x'))$.

Definition 3.2 Given an augmented tLDBA \bar{B}_φ and an MDP M , a tuple $M \otimes \bar{B}_\varphi = M^\otimes = (S^\otimes, A^\otimes, \mathcal{A}^\otimes, s_{init}^\otimes, P^\otimes, \delta^\otimes, \mathcal{F}^\otimes)$ is a product MDP, where $S^\otimes = S \times \bar{X}$ is the finite set of states, $A^\otimes = A$ is the finite set of actions, $\mathcal{A}^\otimes : S^\otimes \rightarrow 2^{A^\otimes}$ is the mapping defined as $\mathcal{A}^\otimes((s, \bar{x})) = \mathcal{A}(s)$, $s_{init}^\otimes = (s_{init}, \bar{x}_{init})$ is the initial states, $P^\otimes : S^\otimes \times S^\otimes \times A^\otimes \rightarrow [0, 1]$ is the transition probability defined as

$$P^\otimes(s^\otimes | s^\otimes, a) = \begin{cases} P(s' | s, a) & \text{if } (\bar{x}, L((s, a, s')), \bar{x}') \in \bar{\delta}, \\ 0 & \text{otherwise,} \end{cases}$$

$\delta^\otimes = \{(s^\otimes, a, s^\otimes') \in S^\otimes \times A^\otimes \times S^\otimes ; P^\otimes(s^\otimes' | s^\otimes, a) > 0\}$ is the set of transitions, and $\mathcal{F}^\otimes = \{\bar{F}_1^\otimes, \dots, \bar{F}_n^\otimes\}$ is the acceptance condition, where $\bar{F}_i^\otimes = \{((s, \bar{x}), a, (s', \bar{x}')) \in \delta^\otimes ; (\bar{x}, L(s, a, s'), \bar{x}') \in \bar{F}_i\}$ for each $i \in \{1, \dots, n\}$.

Definition 3.3 The reward function $\mathcal{R} : S^\otimes \times A^\otimes \times S^\otimes \rightarrow \mathbb{R}_{\geq 0}$ is defined as

$$\mathcal{R}(s^\otimes, a, s^\otimes') = \begin{cases} r_p & \text{if } \exists i \in \{1, \dots, n\}, (s^\otimes, a, s^\otimes') \in \bar{F}_i^\otimes, \\ 0 & \text{otherwise,} \end{cases}$$

where r_p is a positive value.

Under the product MDP M^\otimes and the reward function \mathcal{R} , which is based on the acceptance condition of M^\otimes , we show that if there exists a positional policy π satisfying the LTL specification φ , maximizing the expected discounted reward produces a policy satisfying φ .

For a Markov chain MC_π^\otimes induced by a product MDP M^\otimes with a positional policy π , let $S_\pi^\otimes = T_\pi^\otimes \sqcup R_\pi^{\otimes 1} \sqcup \dots \sqcup R_\pi^{\otimes h}$ be the set of states in MC_π^\otimes , where T_π^\otimes is the set of transient states and $R_\pi^{\otimes i}$ is the recurrent class for each $i \in \{1, \dots, h\}$, and let $R(MC_\pi^\otimes)$ be the set of all recurrent classes in MC_π^\otimes . Let $\delta_{\pi, i}^\otimes$ be the set of transtions in a recurrent class $R_\pi^{\otimes i}$, namely $\delta_{\pi, i}^\otimes = \{(s^\otimes, a, s^\otimes') \in \delta^\otimes ; s^\otimes \in R_\pi^{\otimes i}, P^\otimes(s^\otimes' | s^\otimes, a) > 0\}$, and let $P_\pi^\otimes : S_\pi^\otimes \times S_\pi^\otimes \rightarrow [0, 1]$ be the transition probability under π .

Lemma 3.1 For any policy π and any recurrent class $R_\pi^{\otimes i}$ in the Markov chain MC_π^\otimes , MC_π^\otimes satisfies one of the following conditions.

1. $\delta_{\pi, i}^\otimes \cap \bar{F}_j^\otimes \neq \emptyset, \forall j \in \{1, \dots, n\}$,
2. $\delta_{\pi, i}^\otimes \cap \bar{F}_j^\otimes = \emptyset, \forall j \in \{1, \dots, n\}$.

Lemma 3.1 implies that for an LTL formula φ if a path ρ under a policy π does not satisfy φ , then the agent obtains no reward in recurrent classes; otherwise there exists at least one recurrent class where the agent obtains rewards infinitely often.

Theorem 3.1 Let M^\otimes be the product MDP corresponding to an MDP M and an LTL formula φ . If there exists a positional policy satisfying φ , then there exists a discount factor γ^* such that any algorithm that maximizes the expected reward with $\gamma > \gamma^*$ will find a positional policy satisfying φ .

We show the overall procedure of our proposed method in Algorithm 1. We employ Q-learning in Algorithm 1, but any algorithms maximizing the expected discounted reward can be applied to our proposed method.

Algorithm 1 RL-based synthesis of control policy on the MDP with the augmented tLDBA.

Input: LTL formula φ and MDP M

Output: Optimal policy π^* on the product MDP M^\otimes

- 1: Translate φ into tLDBA B_φ .
 - 2: Augment B_φ to \bar{B}_φ .
 - 3: Construct the product MDP M^\otimes of M and \bar{B}_φ .
 - 4: Initialize $Q : S^\otimes \times A^\otimes \rightarrow \mathbb{R}_{\geq 0}$.
 - 5: Initialize episode length T .
 - 6: **while** Q is not converged **do**
 - 7: $s^\otimes \leftarrow (s_{init}, (x_{init}, \mathbf{0}))$.
 - 8: **for** $t = 1$ to T **do**
 - 9: Choose the action a by a policy π .
 - 10: Observe the next state $s^{\otimes'}$.
 - 11: $Q(s^\otimes, a) \leftarrow Q(s^\otimes, a) + \alpha \{ \mathcal{R}(s^\otimes, a, s^{\otimes'}) + \gamma \max_{a'} Q(s^{\otimes'}, a') - Q(s^\otimes, a) \}$
 - 12: $s^\otimes \leftarrow s^{\otimes'}$
 - 13: **end for**
 - 14: **end while**
-

3.2 Example

In this section, we evaluate our proposed method and compare it with an existing work. We consider a path planning problem of a robot in an environment consisting of eight rooms and one corridor as shown in Fig. 3.1. The state s_7 is the initial state and the action space is specified with $\mathcal{A}(s) = \{Right, Left, Up, Down\}$ for any state $s \neq s_4$ and $\mathcal{A}(s_4) = \{to_s_0, to_s_1, to_s_2, to_s_3, to_s_5, to_s_6, to_s_7, to_s_8\}$, where to_s_i means attempting to go to the state s_i for $i \in \{0, 1, 2, 3, 5, 6, 7, 8\}$. The robot moves in the intended direction with probability 0.9 and it stays in the same state with probability 0.1 if it is in the state s_4 . In the states other than s_4 , it moves in the intended direction with probability 0.9 and it moves in the opposite direction to that it intended to go with probability 0.1. If the robot tries to go to outside the environment, it stays in the same




Fig. 3.1 The environment consisting of eight rooms and one corridor. Red arcs are the transitions that we want to occur infinitely often, while blue arcs are the transitions that we never want to occur. s_7 is the initial state.




Fig. 3.2 The tLDBA recognizing the LTL formula $\mathbf{GF}a \wedge \mathbf{GF}b \wedge \mathbf{G}\neg c$, where the initial state is x_0 . Red arcs are accepting transitions that are numbered in accordance with the accepting sets they belong to, e.g., $\text{① } a \wedge \neg b \wedge \neg c$ means the transition labeled by it belongs to the accepting set F_1 .

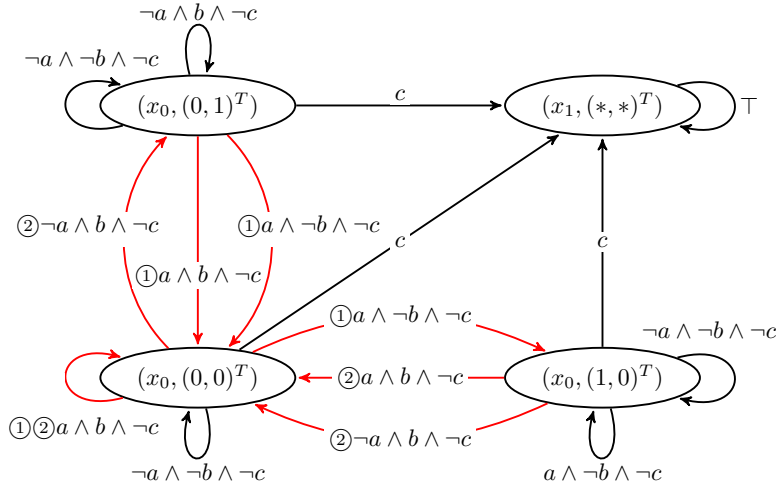


Fig. 3.3 The augmented automaton for the tLDBA in Fig. 3.2 recognizing the LTL formula $\mathbf{GF}a \wedge \mathbf{GF}b \wedge \mathbf{G}\neg c$, where the initial state is $(x_0, (0, 0)^T)$. Red arcs are accepting transitions that are numbered in accordance with the accepting sets they belong to. All states corresponding to x_1 are merged into $(x_1, (*, *)^T)$.

state. The labeling function is as follows.

$$L((s, a, s')) = \begin{cases} \{c\} & \text{if } s' = s_i, i \in \{2, 3, 5, 6\}, \\ \{a\} & \text{if } (s, a, s') = (s_4, to_s_0, s_0), \\ \{b\} & \text{if } (s, a, s') = (s_4, to_s_8, s_8), \\ \emptyset & \text{otherwise.} \end{cases}$$

In the example, the robot tries to take two transitions that we want to occur infinitely often, represented by arcs labeled by $\{a\}$ and $\{b\}$, while avoiding unsafe transitions represented by the arcs labeled by $\{c\}$. This is formally specified by the following LTL formula.

$$\varphi = \mathbf{GF}a \wedge \mathbf{GF}b \wedge \mathbf{G}\neg c.$$

The above LTL formula requires the robot to keep on entering the two rooms s_0 and s_8 from the corridor s_4 regardless of the order of entries, while avoiding entering the four rooms s_2 , s_3 , s_5 , and s_6 .

We use Owl [22] to obtain the tLDBA corresponding to the LTL formula. The tLDBA $B_\varphi = (X, x_{init}, \Sigma, \delta, \mathcal{F})$ and its augmented automaton $\bar{B}_\varphi = (\bar{X}, \bar{x}_{init}, \bar{\Sigma}, \bar{\delta}, \bar{\mathcal{F}})$ are shown in Figs. 3.2 and 3.3, respectively. Specifically, the acceptance condition \mathcal{F} of the tLDBA is given by $\mathcal{F} = \{F_1, F_2\}$, where $F_1 = \{(x_0, \{a\}, x_0), (x_0, \{a, b\}, x_0)\}$ and $F_2 = \{(x_0, \{b\}, x_0), (x_0, \{a, b\}, x_0)\}$.

We use Q-learning with ε -greedy policy and gradually reduce ε to 0 to learn an optimal policy asymptotically. We set the positive reward $r_p = 2$, the epsilon greedy parameter $\varepsilon = \frac{0.95}{n_t(s^\otimes)}$, where $n_t(s^\otimes)$ is the number of visits to state s^\otimes within t time steps [21], and the discount factor $\gamma = 0.9$. The learning rate α varies in accordance with *the Robbins-Monro condition*.

We also evaluate the method by Hasanbeig *et al.*[14] with the same example. They use state-based LDBAs for LTL formulas and construct the product MDP of an MDP and a state-based LDBA to synthesize a policy satisfying the LTL formula. They proposed the accepting frontier function $Acc : X \times 2^X \rightarrow 2^X$ where X is the set of states of the state-based LDBA. Under initializing a set of states \mathbb{F} with the union of the all accepting sets of the state-based LDBA, the function receives the state x after each transition and the set \mathbb{F} . If x is in \mathbb{F} , then Acc removes the accepting sets containing x from \mathbb{F} . The reward function is based on the varying set \mathbb{F} . We conduct the same example with their method using the tLDBA instead.

Figs. 3.4 and 3.5 show the average reward and the optimal policy, respectively, as a result of the learning when using our proposed method and the method in [14] after 10000 iterations and 1000 episodes. The arithmetic mean of average reward in each episode for 20 learning sessions is displayed per 100 episodes in Fig. 3.4.




Fig. 3.4 The arithmetic mean of average reward in each episode for 20 learning sessions obtained from our proposed method (left) and the method by Hasanbeig *et al.*[14] (right). They are plotted per 100 episodes and the green areas represent the range of standard deviations.




Fig. 3.5 The optimal policy obtained from our proposed method (left) and the method by Hasanbeig *et al.*[14] (right).

The results suggest that our proposed method can synthesize a policy satisfying φ on the MDP, while the method in [14] cannot. This is because it is impossible that the transitions labeled by $\{a\}$ and $\{b\}$ occur from s_4 infinitely often by any positional policy with the tLDBA. In detail, the state of the tLDBA is always x_0 while the agent does not move to states s_2 , s_3 , s_5 , and s_6 . Thus, the state of the product MDP is always (s_4, x_0) while the agent stays in s_4 . Therefore, the method in [14] may not synthesize policies satisfying LTL specifications depending on the setting of MDPs or LTL specifications.

Chapter 4

4th Chapter

Chapter 5

Conclusions

A

Proofs

Proof (Proof of lemma 3.1) Suppose that MC_π^\otimes satisfies neither conditions 1 nor 2. Then, there exists a policy π , $i \in \{1, \dots, h\}$, and $j_1, j_2 \in \{1, \dots, n\}$ such that $\delta_{\pi,i}^\otimes \cap \bar{F}_{j_1}^\otimes = \emptyset$ and $\delta_{\pi,i}^\otimes \cap \bar{F}_{j_2}^\otimes \neq \emptyset$. In other words, there exists a nonempty and proper subset $J \in 2^{\{1, \dots, n\}} \setminus \{\{1, \dots, n\}, \emptyset\}$ such that $\delta_{\pi,i}^\otimes \cap \bar{F}_j^\otimes \neq \emptyset$ for any $j \in J$. For any transition $(s^\otimes, a, s^{\otimes'}) \in \delta_{\pi,i}^\otimes \cap \bar{F}_j^\otimes$, the following equation holds by the properties of the recurrent states in MC_π^\otimes [18].

$$\sum_{k=0}^{\infty} p^k((s^\otimes, a, s^{\otimes'}), (s^\otimes, a, s^{\otimes'})) = \infty, \quad (\text{A.1})$$

where $p^k((s^\otimes, a, s^{\otimes'}), (s^\otimes, a, s^{\otimes'}))$ is the probability that the transition $(s^\otimes, a, s^{\otimes'})$ occurs again after the occurrence of itself in k time steps. Eq. (A.1) means that the agent obtains a reward infinitely often. This contradicts the definition of the acceptance condition of the product MDP M^\otimes . \square

Proof Suppose that π^* is an optimal policy but does not satisfy the LTL formula φ . Then, for any recurrent class $R_{\pi^*}^{\otimes i}$ in the Markov chain $MC_{\pi^*}^\otimes$ and any accepting set \bar{F}_j^\otimes of the product MDP M^\otimes , $\delta_{\pi^*,i}^\otimes \cap \bar{F}_j^\otimes = \emptyset$ holds by Lemma 3.1. Thus, the agent under the policy π^* can obtain rewards only in the set of transient states. We consider the best scenario in the assumption. Let $p^k(s, s')$ be the probability of going to a state s' in k time steps after leaving the state s , and let $Post(T_{\pi^*}^\otimes)$ be the set of states in recurrent classes that can be transitioned from states in $T_{\pi^*}^\otimes$ by one action. For the initial state s_{init}^\otimes in the set of transient states, it holds that

$$\begin{aligned} V^{\pi^*}(s_{init}^\otimes) &= \sum_{k=0}^{\infty} \sum_{s^\otimes \in T_{\pi^*}^\otimes} \gamma^k p^k(s_{init}^\otimes, s^\otimes) \sum_{s^{\otimes'} \in T_{\pi^*}^\otimes \cup Post(T_{\pi^*}^\otimes)} P_{\pi^*}^\otimes(s^{\otimes'} | s^\otimes) \mathcal{R}(s^\otimes, a, s^{\otimes'}) \\ &\leq r_p \sum_{k=0}^{\infty} \sum_{s^\otimes \in T_{\pi^*}^\otimes} \gamma^k p^k(s_{init}^\otimes, s^\otimes), \end{aligned}$$

where the action a is selected by π^* . By the property of the transient states, for any state s^\otimes in $T_{\pi^*}^\otimes$, there exists a bounded positive value m such that $\sum_{k=0}^{\infty} \gamma^k p^k(s_{init}^\otimes, s^\otimes) \leq \sum_{k=0}^{\infty} p^k(s_{init}^\otimes, s^\otimes) < m$ [18]. Therefore, there exists a bounded positive value \bar{m} such that $V^{\pi^*}(s_{init}^\otimes) < \bar{m}$. Let $\bar{\pi}$ be a positional policy satisfying φ . We consider the following two cases.

1. Assume that the initial state s_{init}^\otimes is in a recurrent class $R_{\bar{\pi}}^{\otimes i}$ for some $i \in \{1, \dots, h\}$. For any accepting set \bar{F}_j^\otimes , $\delta_{\bar{\pi}, i}^\otimes \cap \bar{F}_j^\otimes \neq \emptyset$ holds by the definition of $\bar{\pi}$. The expected discounted reward for s_{init}^\otimes is given by

$$V^{\bar{\pi}}(s_{init}^\otimes) = \sum_{k=0}^{\infty} \sum_{s^\otimes \in R_{\bar{\pi}}^{\otimes i}} \gamma^k p^k(s_{init}^\otimes, s^\otimes) \sum_{s^{\otimes'} \in R_{\bar{\pi}}^{\otimes i}} P_{\bar{\pi}}^\otimes(s^{\otimes'} | s^\otimes) \mathcal{R}(s^\otimes, a, s^{\otimes'}),$$

where the action a is selected by $\bar{\pi}$. Since s_{init}^\otimes is in $R_{\bar{\pi}}^{\otimes i}$, there exists a positive number $\bar{k} = \min\{k; k \geq n, p^k(s_{init}^\otimes, s_{init}^\otimes) > 0\}$ [18]. We consider the worst scenario in this case. It holds that

$$\begin{aligned} V^{\bar{\pi}}(s_{init}^\otimes) &\geq \sum_{k=n}^{\infty} p^k(s_{init}^\otimes, s_{init}^\otimes) (\gamma^k + \gamma^{k-1} + \dots + \gamma^{k-n+1}) r_p \\ &\geq \sum_{k=1}^{\infty} p^{k\bar{k}}(s_{init}^\otimes, s_{init}^\otimes) (\gamma^{k\bar{k}} + \dots + \gamma^{k\bar{k}-n+1}) r_p \\ &> r_p \sum_{k=1}^{\infty} \gamma^{k\bar{k}} p^{k\bar{k}}(s_{init}^\otimes, s_{init}^\otimes), \end{aligned}$$

whereas all states in $R(MC_{\bar{\pi}}^\otimes)$ are positive recurrent because $|S^\otimes| < \infty$ [19]. Obviously, $p^{k\bar{k}}(s_{init}^\otimes, s_{init}^\otimes) \geq (p^{\bar{k}}(s_{init}^\otimes, s_{init}^\otimes))^k > 0$ holds for any $k \in (0, \infty)$ by the Chapman-Kolmogorov equation [18]. Furthermore, we have $\lim_{k \rightarrow \infty} p^{k\bar{k}}(s_{init}^\otimes, s_{init}^\otimes) > 0$ by the property of irreducibility and positive recurrence [20]. Hence, there exists \bar{p} such that $0 < \bar{p} < p^{k\bar{k}}(s_{init}^\otimes, s_{init}^\otimes)$ for any $k \in (0, \infty]$ and we have

$$V^{\bar{\pi}}(s_{init}^\otimes) > r_p \bar{p} \gamma^{\bar{k}} p^{\bar{k}}(s_{init}^\otimes, s_{init}^\otimes) \frac{1}{1 - \gamma^{\bar{k}}}.$$

Therefore, for any $\bar{m} \in (V^{\pi^*}(s_{init}^\otimes), \infty)$ and any $r_p < \infty$, there exists $\gamma^* < 1$ such that $\gamma > \gamma^*$ implies $V^{\bar{\pi}}(s_{init}^\otimes) > r_p \bar{p} \gamma^{\bar{k}} p^{\bar{k}}(s_{init}^\otimes, s_{init}^\otimes) \frac{1}{1 - \gamma^{\bar{k}}} > \bar{m}$.

2. Assume that the initial state s_{init}^\otimes is in the set of transient states $T_{\bar{\pi}}^\otimes$. $P_{\bar{\pi}}^{M^\otimes}(s_{init}^\otimes \models \varphi) > 0$ holds by the definition of $\bar{\pi}$. For a recurrent class $R_{\bar{\pi}}^{\otimes i}$ such that $\delta_{\bar{\pi}, i}^\otimes \cap \bar{F}_j^\otimes \neq \emptyset$ for each accepting set \bar{F}_j^\otimes , there exist a number $\bar{l} > 0$, a state \hat{s}^\otimes in $Post(T_{\bar{\pi}}^\otimes) \cap R_{\bar{\pi}}^{\otimes i}$, and a subset of transient states $\{s_1^\otimes, \dots, s_{\bar{l}-1}^\otimes\} \subset T_{\bar{\pi}}^\otimes$ such that $p(s_{init}^\otimes, s_1^\otimes) > 0$,

$p(s_i^\otimes, s_{i+1}^\otimes) > 0$ for $i \in \{1, \dots, \bar{l}-2\}$, and $p(s_{\bar{l}-1}^\otimes, \hat{s}^\otimes) > 0$ by the property of transient states. Hence, it holds that $p^{\bar{l}}(s_{init}^\otimes, \hat{s}^\otimes) > 0$ for the state \hat{s}^\otimes . Thus, by ignoring rewards in T_π^\otimes , we have

$$\begin{aligned} V^{\bar{\pi}}(s_{init}^\otimes) &\geq P_\pi^{M^\otimes}(s_{init}^\otimes \models \varphi) \gamma^{\bar{l}} p^{\bar{l}}(s_{init}^\otimes, \hat{s}^\otimes) \sum_{k=0}^{\infty} \sum_{s^{\otimes'} \in R_\pi^{\otimes i}} \gamma^k p^k(\hat{s}^\otimes, s^{\otimes'}) \\ &\quad \sum_{s^{\otimes''} \in R_\pi^{\otimes i}} P_\pi^\otimes(s^{\otimes''} | s^{\otimes'}) \mathcal{R}(s^{\otimes'}, a, s^{\otimes''}) \\ &> P_\pi^{M^\otimes}(s_{init}^\otimes \models \varphi) \gamma^{\bar{l}} p^{\bar{l}}(s_{init}^\otimes, \hat{s}^\otimes) r_p \bar{p} \gamma^{\bar{k}'} p^{\bar{k}'}(\hat{s}^\otimes, \hat{s}^\otimes) \frac{1}{1 - \gamma^{\bar{k}'}} \end{aligned}$$

where $\bar{k}' \geq n$ is a constant and $0 < \bar{p} < p^{k\bar{k}'}(\hat{s}^\otimes, \hat{s}^\otimes)$ for any $k \in (0, \infty]$. Therefore, for any $\bar{m} \in (V^{\pi^*}(s_{init}^\otimes), \infty)$ and any $r_p < \infty$, there exists $\gamma^* < 1$ such that $\gamma > \gamma^*$ implies $V^{\bar{\pi}}(s_{init}^\otimes) > P_\pi^{M^\otimes}(s_{init}^\otimes \models \varphi) \gamma^{\bar{l}} p^{\bar{l}}(s_{init}^\otimes, \hat{s}^\otimes) r_p \bar{p} \gamma^{\bar{k}'} p^{\bar{k}'}(\hat{s}^\otimes, \hat{s}^\otimes) \frac{1}{1 - \gamma^{\bar{k}'}} > \bar{m}$.

The results contradict the optimality assumption of π^* . □

Acknowledgment

I would like to express my deep sense of gratitude to my adviser Professor Toshimitsu Ushio, Graduate School of Engineering Science, Osaka University, for his invaluable, constructive advice and constant encouragement during this work. Professor Ushio's deep knowledge and his eye for detail have inspired me much.

References

- [1] C. Baier and J.-P. Katoen, *Principles of Model Checking*. MIT Press, 2008.
- [2] E. M. Clarke, Jr., O. Grumberg, D. Kroening, D. Peled, and H. Veith, *Model Checking*, 2nd Edition. MIT Press, 2018.
- [3] M. Kloetzer, C. Belta, “A fully automated framework for control of linear systems from temporal logic specifications,” *IEEE Trans. Autom. Contr.*, vol. 53, no. 1, pp. 287–297, 2008.
- [4] H. Kress-Gazit, G. E. Fainekos, and G. J. Pappas, “Temporal-logic-based reactive mission and motion planning,” *IEEE Trans. Robotics*, vol. 25, no. 6, pp. 1370–1381, 2009.
- [5] T. Wongpiromsarn, U. Topcu, and R. M. Murray, “Receding horizon temporal logic planning,” *IEEE Trans. Autom. Contr.*, vol. 57, no. 11, pp. 2817–2830, 2012.
- [6] A. Sakakibara and T. Ushio, “Decentralized supervision and coordination of concurrent discrete event systems under LTL constraints,” in *Proc. 14th International Workshop on Discrete Event Systems*, 2018, pp. 18–23.
- [7] C. Belta, B. Yordanov, and E. A. Gol, *Formal Methods for Discrete-Time Dynamical Systems*. Springer, 2017.
- [8] M. L. Puterman, *Markov Decision Processes, Discrete Stochastic Dynamic Programming*. John Wiley & Sons, Inc., 1994.
- [9] E. M. Wolff, U. Topcu, and R. M. Murray, “Robust control of uncertain Markov decision processes with temporal logic specifications,” in *Proc. 51st IEEE Conference on Decision and Control*, 2012, pp. 3372–3379.
- [10] D. Sadigh, E. S. Kim, A. Coogan, S. S. Sastry, and S. Seshia, “A learning based approach to control synthesis of Markov decision processes for linear temporal logic specifications,” in *Proc. 53rd IEEE Conference on Decision and Control*, pp. 1091–1096, 2014.
- [11] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd Edition. MIT Press, 2018.
- [12] M. Hiromoto and T. Ushio, “Learning an optimal control policy for a Markov decision process under linear temporal logic specifications,” in *Proc. 2015 IEEE Symposium on Adaptive Dynamic Programming and Reinforcement Learning*, 2015, pp. 548–555.
- [13] S. Sickert, J. Esparaza, S. Jaax, and J. Křetínský, “Limit-deterministic Büchi au-

- tomata for linear temporal logic,” in *International Conference on Computer Aided Verification*, 2016, pp. 312–332.
- [14] M. Hasanbeig, A. Abate, and D. Kroening, “Logically-constrained reinforcement learning,” *arXiv:1801.08099v8*, Feb. 2019.
 - [15] E. M. Hahn, M. Perez, S. Schewe, F. Somenzi, A. Triverdi, and D. Wojtczak, “Omega-regular objective in model-free reinforcement learning,” *Lecture Notes in Computer Science*, no. 11427, pp. 395–412, 2019.
 - [16] M. Hasanbeig, Y. Kantaros, A. Abate, D. Kroening, G. J. Pappas, and I. Lee, “Reinforcement learning for temporal logic control synthesis with probabilistic satisfaction guarantee,” *arXiv:1909.05304v1*, 2019.
 - [17] A. K. Bozkurt, Y. Wang, M. Zavlanos, and M. Pajic, “Control synthesis from linear temporal logic specifications using model-free reinforcement learning,” *arXiv:1909.07299*, 2019.
 - [18] R. Durrett, *Essentials of Stochastic Processes*, 2nd Edition. ser. Springer texts in statistics. New York; London; Springer, 2012.
 - [19] L. Breuer, “Introduction to Stochastic Processes”, [Online]. Available: <https://www.kent.ac.uk/smsas/personal/lb209/files/sp07.pdf>
 - [20] S.M. Ross, *Stochastic Processes*, 2nd Edition. University of California, Wiley, 1995.
 - [21] S. Singh, T. Jaakkola, M. L. Littman, and C. Szepesvári, “Convergence results for single-step on-policy reinforcement learning algorithms” *Machine Learning*, vol. 38, no. 3, pp. 287–308, 1998.
 - [22] J. Kretínský, T. Meggendorfer, S. Sickert, “Owl: A library for ω -words, automata, and LTL,” in *Proc. 16th International Symposium on Automated Technology for Verification and Analysis*, 2018, pp. 543–550.